New Project Onboarding – AI Foundry

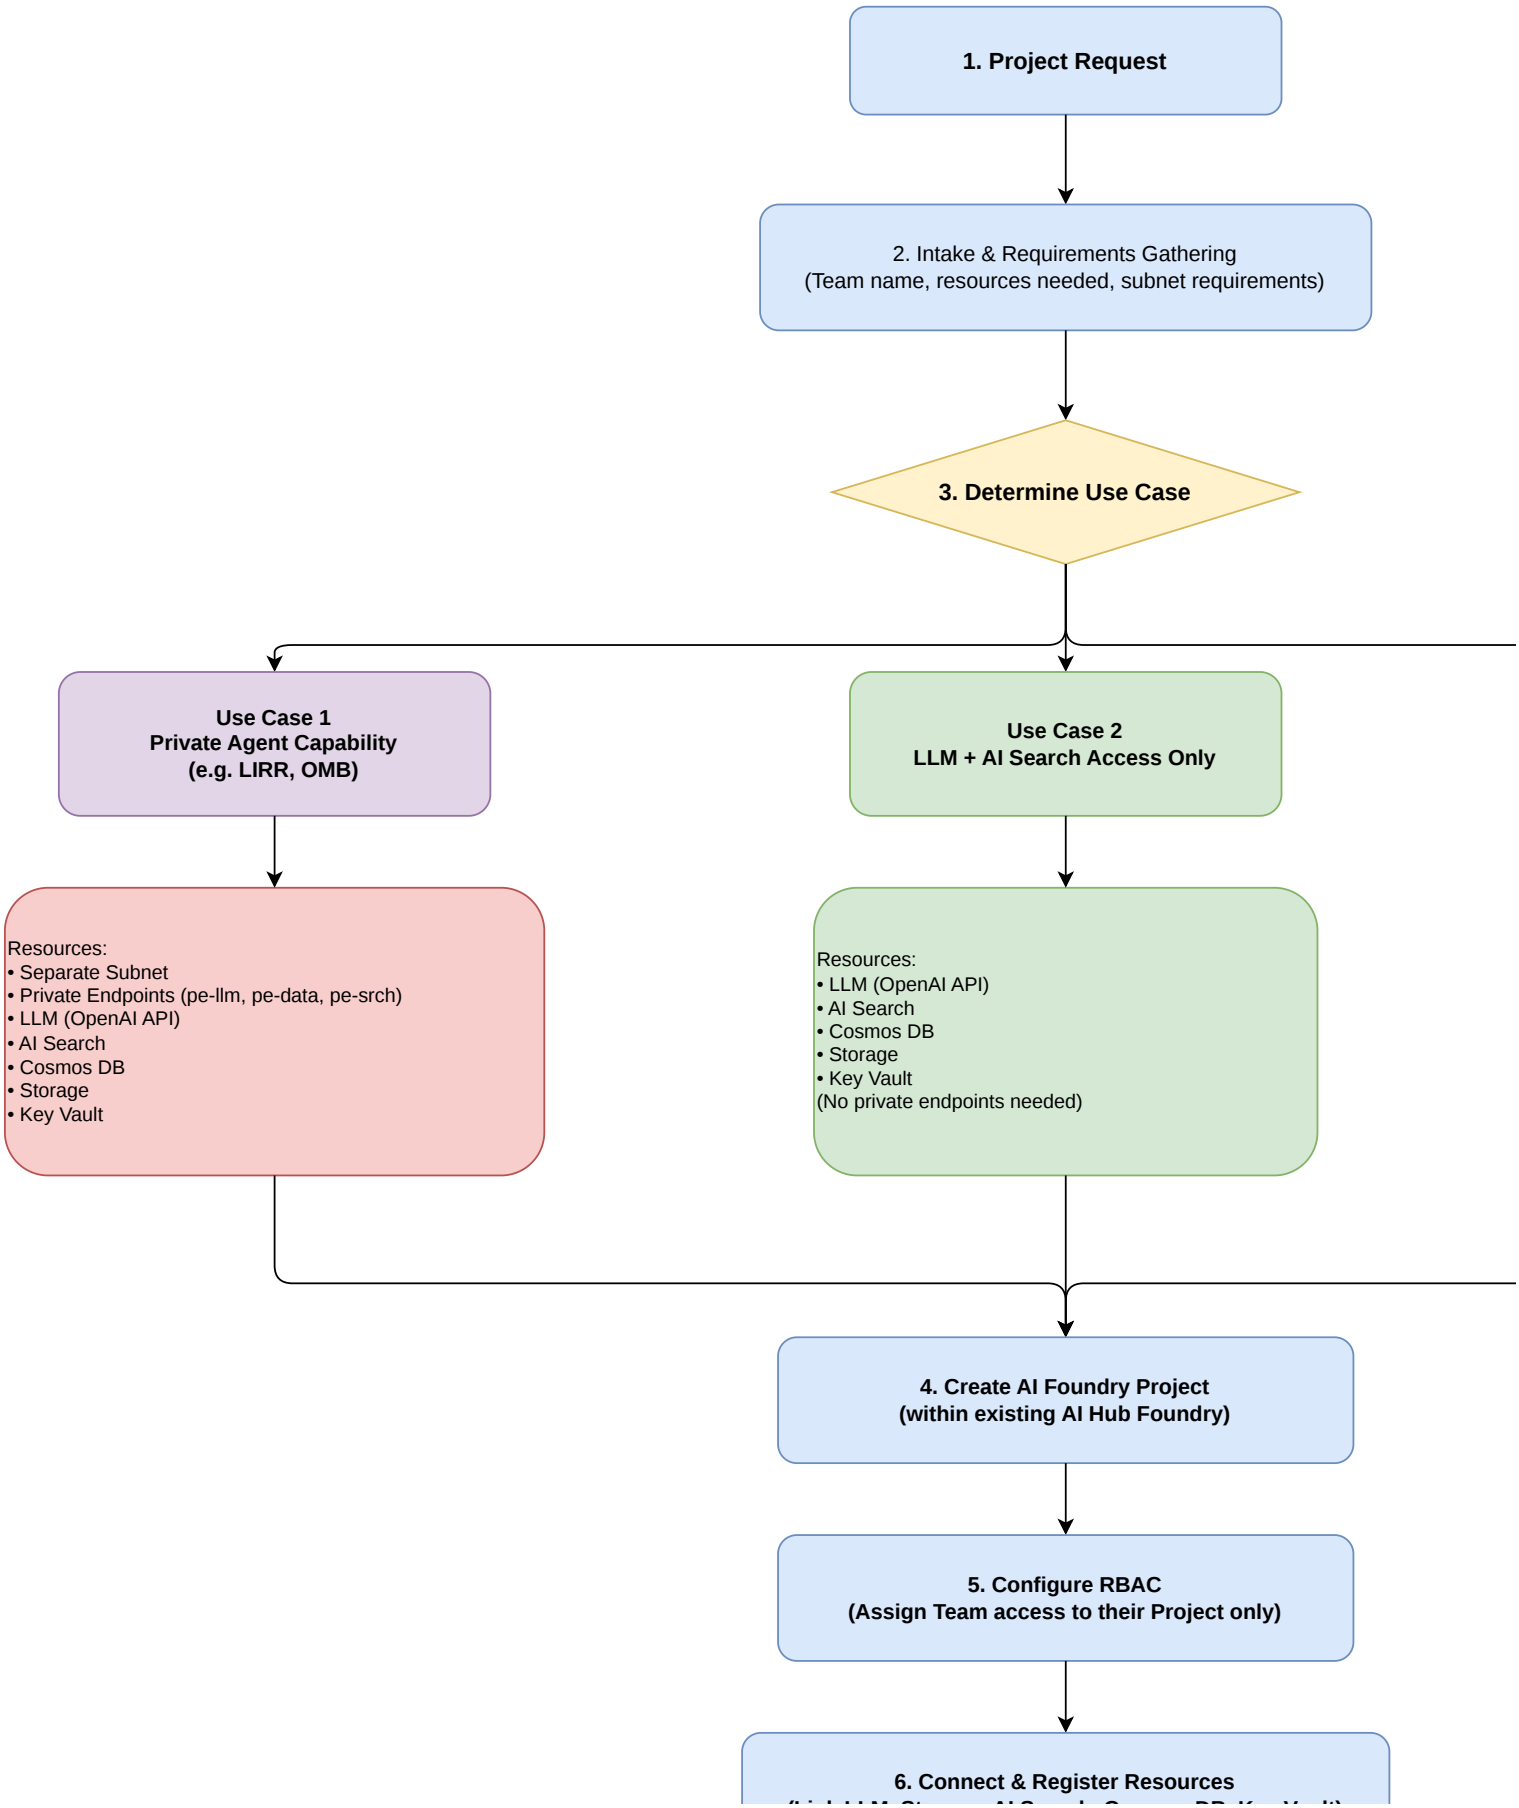

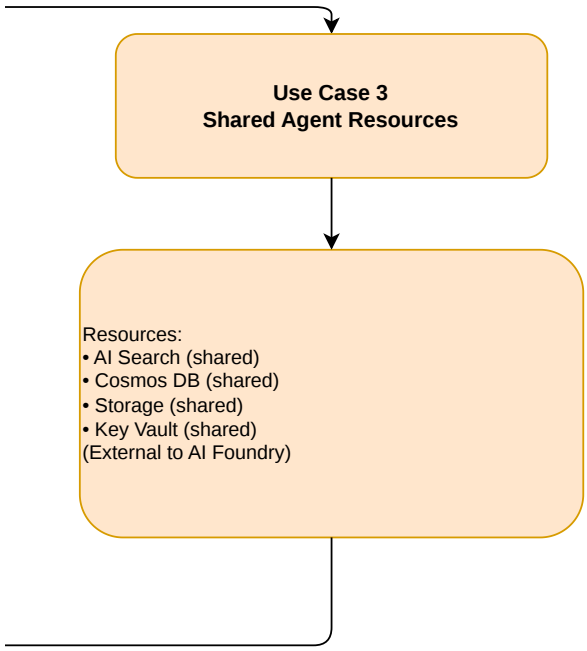
Use Case 3
Shared Agent Resources
Resources:
• AI Search (shared)
• Cosmos DB (shared)
• Storage (shared)
• Key Vault (shared)
(External to AI Foundry)

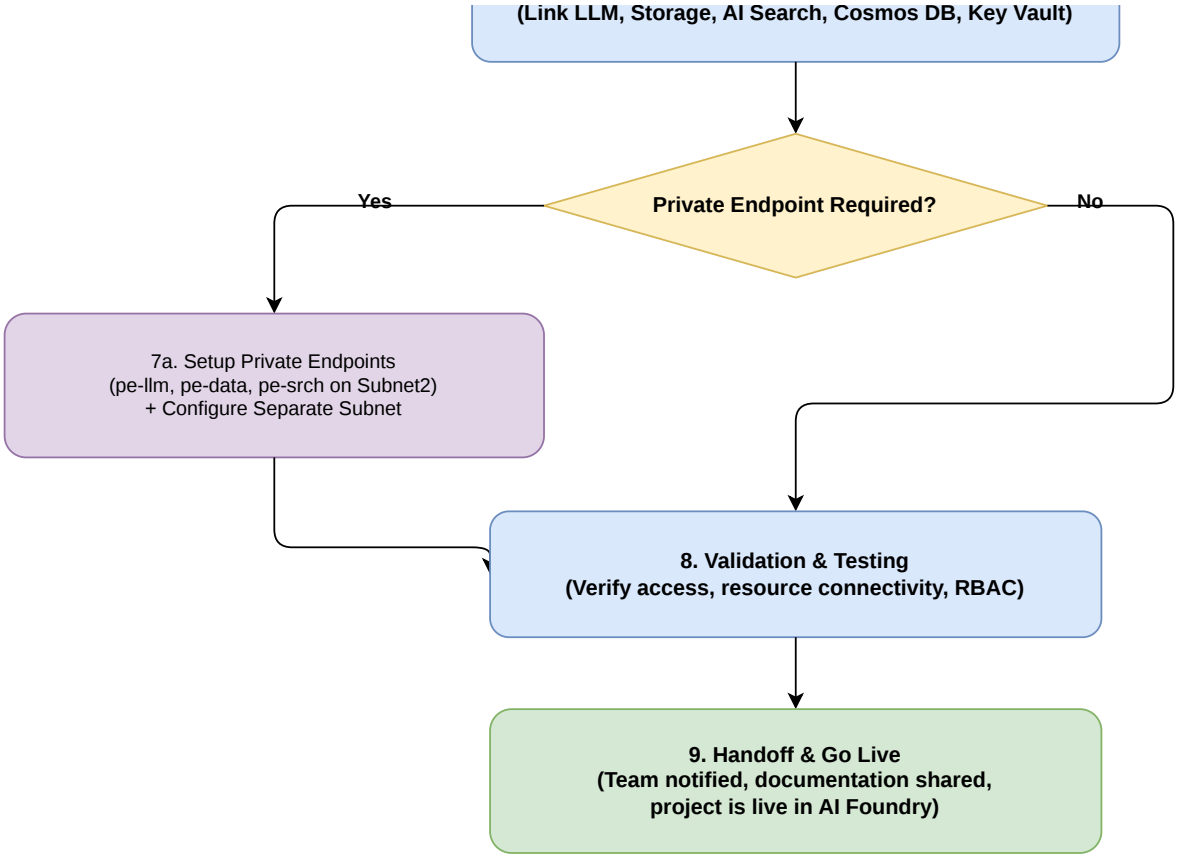
(Link LLM, Storage, AI Search, Cosmos DB, Key Vault)
Private Endpoint Required?
Yes
No
7a. Setup Private Endpoints
(pe-llm, pe-data, pe-srch on Subnet2)
+ Configure Separate Subnet
8. Validation & Testing
(Verify access, resource connectivity, RBAC)
9. Handoff & Go Live
(Team notified, documentation shared,
project is live in AI Foundry)

Azure Landing Zone – AI F

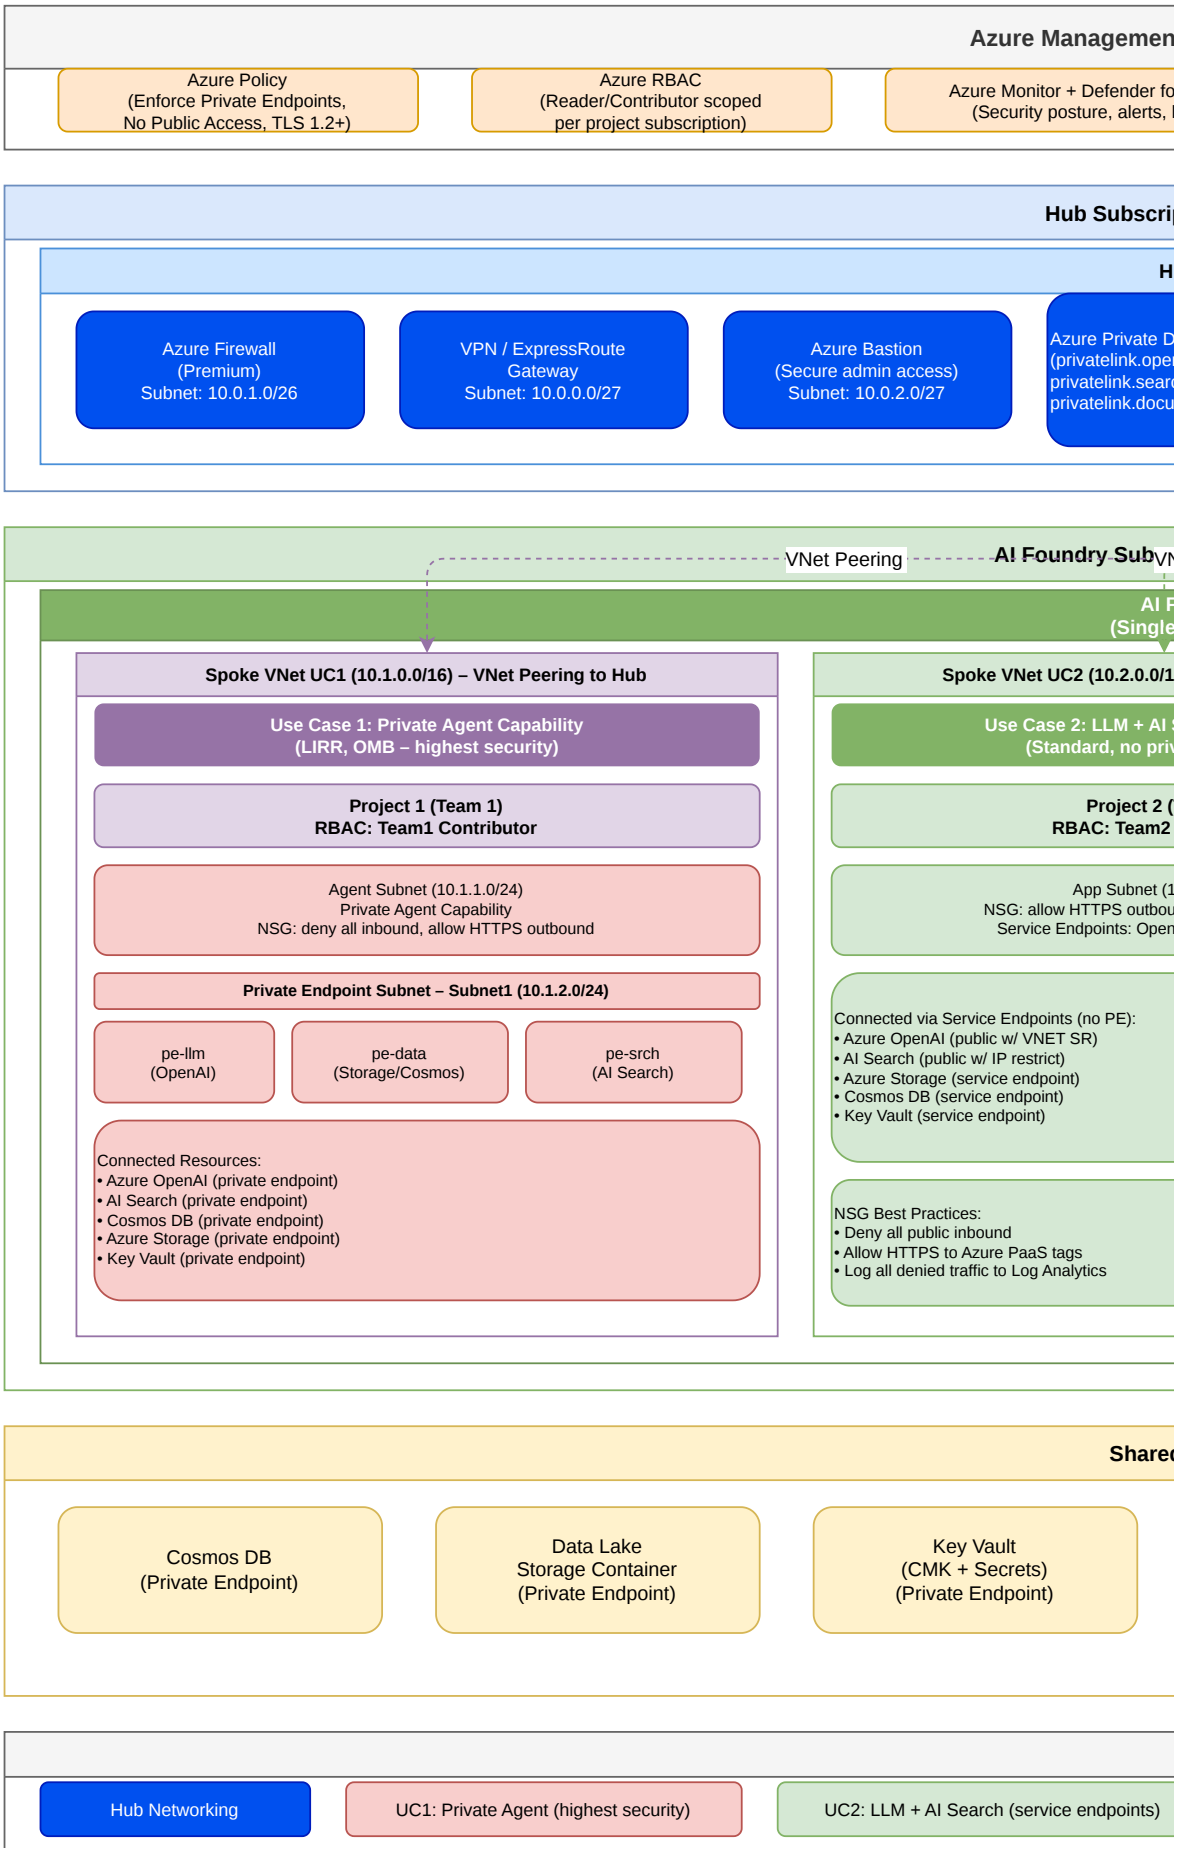

Foundry Best Practices (3 Use Cases)

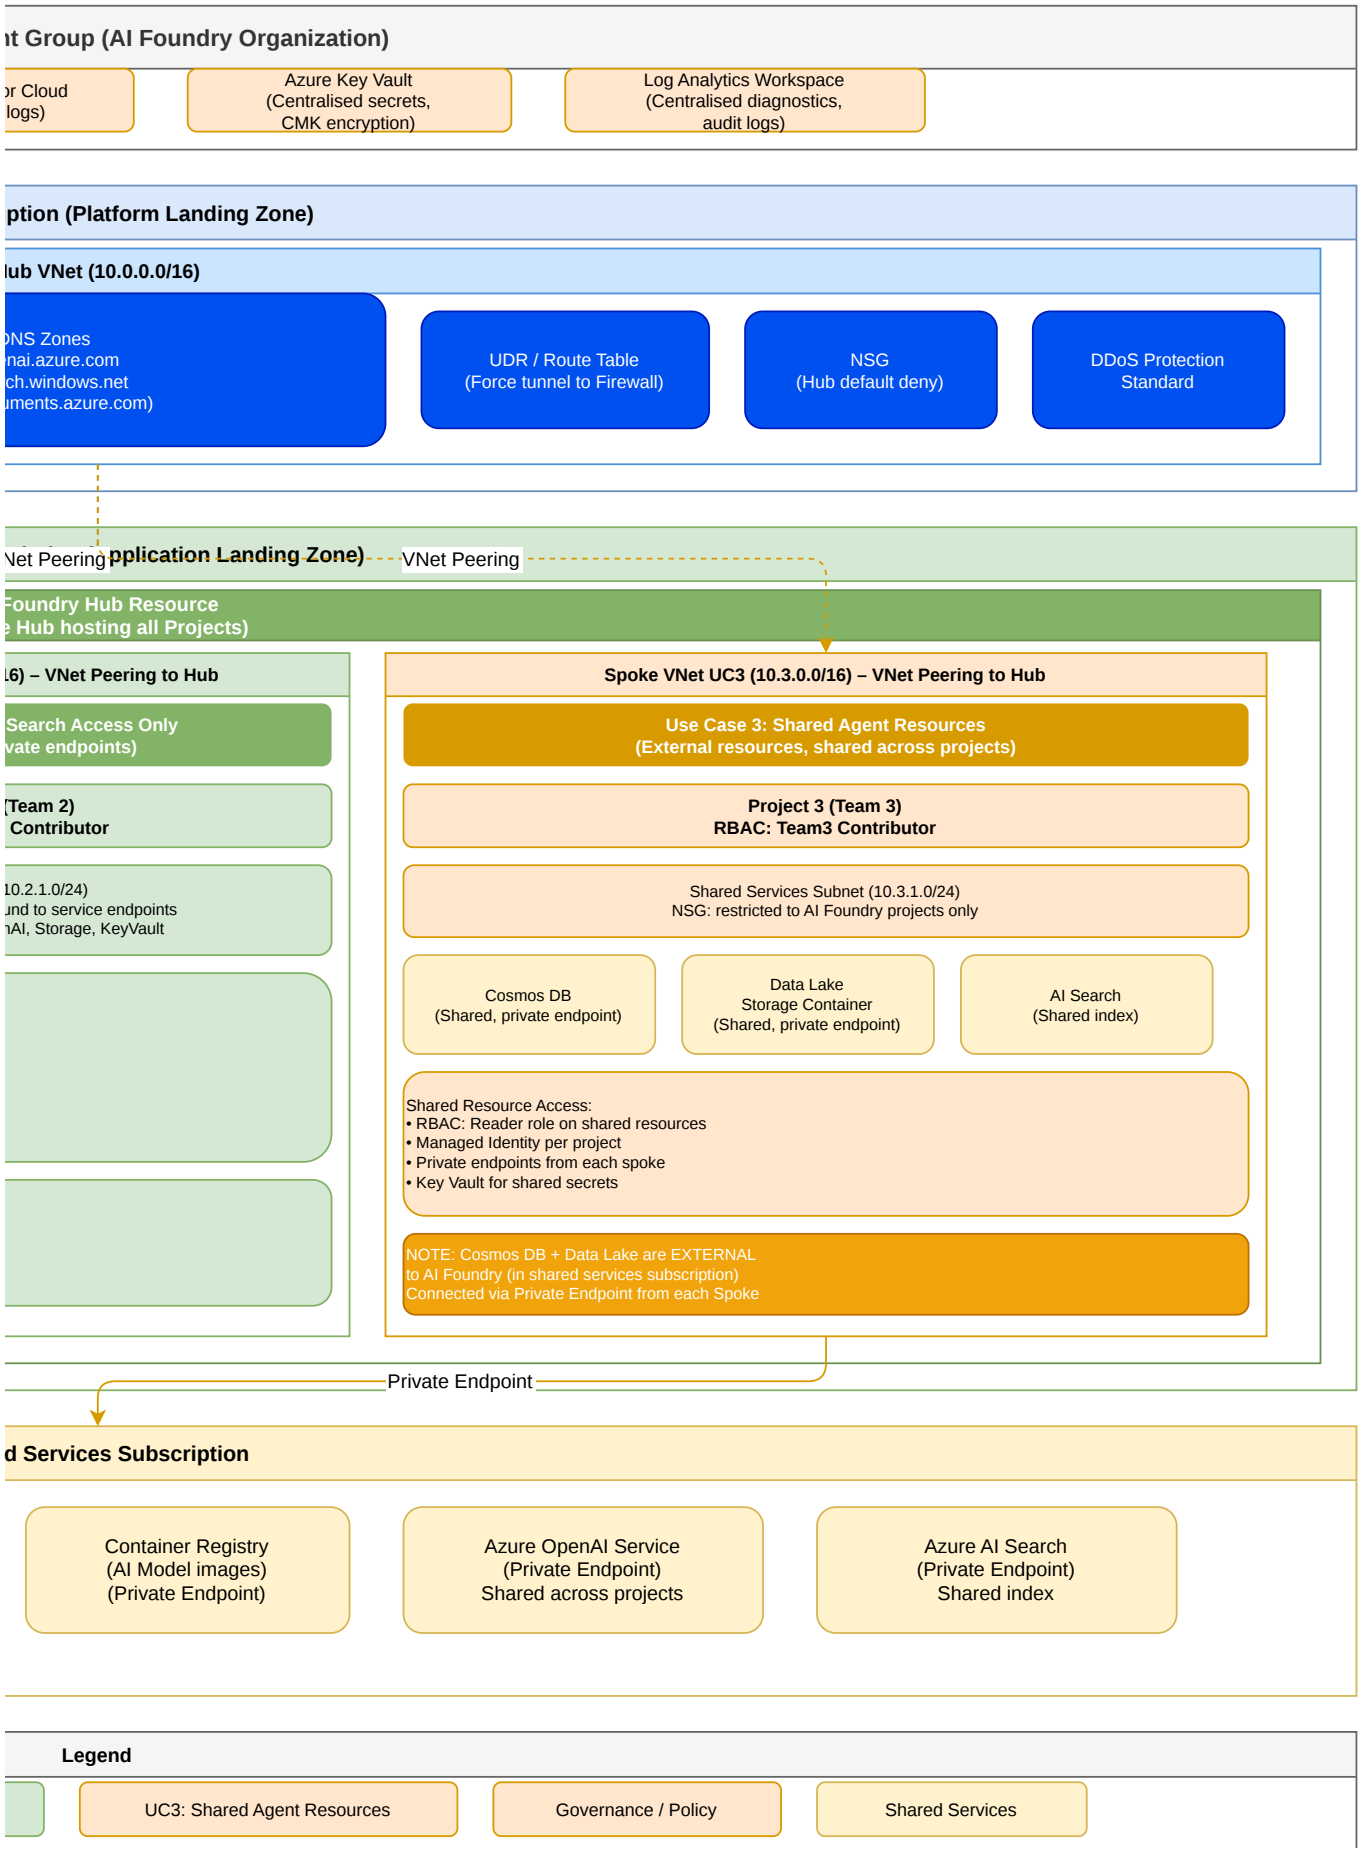

AI Foundry Platform – PRD & Technical Design Document

Version: 1.0 | Status: DRAFT | Date: March 2026 | Owner: AI Platform Team

1. Executive Summary

This document defines the product requirements and technical design for the **AI Foundry Platform** — a centralised Azure AI Hub that enables multiple agency tea

The platform supports three onboarding patterns (Use Cases) based on security, connectivity, and resource-sharing needs. All projects are hosted within a single A

Key objectives: Reduce time-to-onboard for new AI projects • Enforce security and compliance by default • Enable resource sharing without cross-team data leaka

2. Problem Statement

Agency teams require access to large language models, AI search, and data services to build intelligent applications. Without a shared platform:

- Each team independently provisions Azure AI resources, leading to cost duplication and configuration drift.
- There is no consistent security baseline — some teams expose resources publicly, others lack private endpoint configurations.
- Teams with strict compliance needs (e.g. LIRR, OMB requiring private agent capability) have no templated pattern to follow.
- Onboarding takes weeks due to ad-hoc infrastructure requests with no standard process.

3. Goals & Success Metrics

Goal	Success Metric	
Fast project onboarding	Time from request to live project	< 5 business days
Security compliance	% projects with no public endpoint exposure	100% for UC1; 0 criti
RBAC isolation	Zero cross-team data access incidents	0 incidents per quarte
Cost efficiency	Shared resource utilisation vs per-team provisioning	30%+ cost reduction

4. Use Cases & Requirements

Use Case 1 – Private Agent Capability (e.g. LIRR, OMB)

Description: Highest security tier. Teams require a fully private AI agent with no public internet exposure. All Azure PaaS resources are accessed exclusively via Private Endpoints
Trigger: Team handles sensitive data (PII, regulated data) or operates in a compliance-constrained environment.
Functional Requirements: (1) Dedicated Agent Subnet with NSG deny-all inbound. (2) Private Endpoint Subnet (Subnet1) hosting pe-llm, pe-data, pe-srch. (3) All resources (Ope
Non-Functional: Data must not traverse public internet. Audit logs to centralised Log Analytics. RBAC: Team Contributor scoped to this project only.

Use Case 2 – LLM + AI Search Access Only (Standard)

Description: Standard tier. Teams need access to LLM (OpenAI) and AI Search but do not require a fully isolated private network. Resources are accessed via VNet Service Endp
Trigger: Team works with non-sensitive or internal data and primarily needs AI inference and search capabilities.
Functional Requirements: (1) App Subnet with outbound HTTPS to Azure PaaS service tags. (2) Azure OpenAI accessed via VNet service restriction. (3) AI Search accessed via
Non-Functional: All denied traffic logged to Log Analytics. NSG flow logs enabled. No public inbound to subnet. RBAC: Team Contributor scoped to this project only.

Use Case 3 – Shared Agent Resources (External to AI Foundry)

Description: Teams need to connect their AI Foundry project to shared enterprise resources (Cosmos DB, Data Lake, AI Search) that exist outside of the AI Foundry Hub, in a sha
Trigger: Team's data already resides in a shared platform (e.g. enterprise data lake) and duplication is not acceptable. Resources are shared across multiple AI Foundry projects.
Functional Requirements: (1) Project connects to external Cosmos DB and Data Lake via Private Endpoint cross-subscription. (2) Shared AI Search index accessible via Manage
Non-Functional: Shared resources must enforce resource-level access control (no wildcard grants). Audit trail for all shared resource access. Teams cannot modify shared resour

ms to onboard AI projects securely and efficiently within a shared infrastructure.

I Foundry Hub, isolated by Azure RBAC and network controls, avoiding the overhead of separate subscriptions per team.

age • Support private, semi-private, and shared resource patterns.

Target
ical findings in Defender
er
vs isolated deployments

within a dedicated subnet.

enAI, AI Search, Cosmos DB, Storage, Key Vault) accessible only via private endpoints. (4) Azure Private DNS zones registered in Hub. (5) No public IP or service endpoint allowed.

oints or IP-restricted public endpoints.

IP allowlist + AAD auth. (4) Cosmos DB and Storage via Service Endpoints. (5) Key Vault via Service Endpoint for secrets.

red services subscription or resource group.

ed Identity. (3) RBAC: Reader role on shared resources per team Managed Identity. (4) Key Vault in shared subscription holds connection strings. (5) Private DNS entries registered in Hub DNS ces — read-only.

zone.

AI Foundry Platform – PRD & Technical Design Document (cont.)

5. Architecture & Design Decisions

Decision	Choice Made	R
Subscription model	Single subscription, multiple projects	Teams isolated by RBAC + VNet/Subnet. No compliance
AI Foundry topology	Single Hub, multiple Projects	Reduces infrastructure footprint. Shared Hub connections
Network isolation	Hub-Spoke VNet with VNet Peering	Centralised egress via Azure Firewall. DNS via Hub Priva
UC1 private access	Private Endpoints only (pe-llm, pe-data, pe-srch)	Eliminates any public internet data path. Mandatory for PI
UC2 access method	Service Endpoints + VNet restrictions	Lower cost than private endpoints, sufficient for non-sens
UC3 external resources	Cross-sub Private Endpoint + Managed Identity	Avoids data duplication. Managed Identity removes crede

6. Security & Networking Controls

Identity & Access: All resources use Azure Active Directory (Entra ID) authentication. No shared keys or connection strings in code. Managed Identity used per project. RBAC sco

Network: Hub VNet (10.0.0.0/16) with Azure Firewall Premium for centralised egress. Each project gets a dedicated Spoke VNet (UC1: 10.1.x, UC2: 10.2.x, UC3: 10.3.x) peered to

Data Protection: All storage encrypted with Customer-Managed Keys (CMK) stored in centralised Key Vault. TLS 1.2+ enforced via Azure Policy. Soft delete and purge protection

Monitoring: All resources send diagnostic logs to centralised Log Analytics Workspace. NSG flow logs enabled. Azure Monitor alerts for unusual data egress, failed auth attempts,

7. Resource Naming Conventions

Resource Type	Naming Pattern	E
Resource Group	rg-aif-{env}-{project}-{region}	rg-aif-prod-lirr-eus
AI Foundry Hub	aih-{org}-{env}-{region}	aih-nymta-prod-eus
AI Foundry Project	aip-{team}-{env}	aip-lirr-prod / aip-omb-prod
VNet (Spoke)	vnet-spoke-{team}-{env}	vnet-spoke-lirr-prod
Private Endpoint	pe-{service}-{team}-{env}	pe-llm-lirr-prod / pe-srch-lirr-prod
Key Vault	kv-{team}-{env}-{region}	kv-lirr-prod-eus (max 24 chars)

8. Onboarding Process (Step-by-Step)

Step 1 – Project Request: Team submits onboarding request via service desk ticket. Required fields: team name, project name, use case (UC1/UC2/UC3), data classification, list
Step 2 – Intake & Requirements: Platform team reviews request, confirms use case, validates subnet CIDR availability, and assigns a VNet address space.
Step 3 – Create AI Foundry Project: Platform team creates a new Project under the shared AI Hub using the appropriate IaC template (Bicep/Terraform module per use case).
Step 4 – Configure RBAC: Assign team AAD group to Contributor role scoped to their AI Foundry Project. Managed Identity created for the project.
Step 5 – Connect Resources: Register required connections (OpenAI, AI Search, Cosmos DB, Storage, Key Vault) to the project. For UC1: deploy private endpoints. For UC2: co
Step 6 – Private Endpoints (UC1 only): Deploy pe-llm, pe-data, pe-srch in Subnet1. Register DNS A-records in Hub Private DNS Zones.
Step 7 – Validation: Platform team runs connectivity tests and verifies RBAC isolation (no cross-project access). Team receives access credentials.
Step 8 – Handoff: Team notified, documentation and runbook shared. Project marked **LIVE**.

9. Risks, Assumptions & Open Questions

Item	Description	Statu
RISK	Single subscription quota limits may block multiple teams provisioning simultaneously	Mitigation: Pre-request quota increases. Monitor usage via
RISK	UC1 private endpoint DNS misconfig could cause connectivity failures silently	Mitigation: Automated DNS validation script in onboarding p
ASSUMPTION	All teams will use AAD-backed identities (no local accounts or service principals with passwords)	To be confirmed with security team. Owner: Security Archite

Rationale
or billing boundary requiring separate subs.
s (OpenAI, Search) benefit all teams.
ate DNS Zones. UDR for force tunnelling.
II/regulated data environments.
sitive workloads. Simpler to manage.
ential management risk.

ped to project level — teams cannot access sibling projects.

o Hub. UDR forces all outbound through Firewall. NSGs deny all inbound by default, with explicit outbound rules only. Private DNS Zones in Hub resolve all privatelink.* domains.

enabled on all Key Vaults. Defender for Cloud enabled across subscription.

and policy violations.

Example

of resources needed, and designated technical lead.

nfigure service endpoints. For UC3: register cross-sub private endpoints.

us / Owner
Azure Cost + Quota alerts. Owner: Platform Team
pipeline. Owner: Network Team
act

ASSUMPTION	All teams will use AAD-backed identities (no local accounts or service principals with passwords)	To be confirmed with security team. Owner: Security Archite
OPEN Q	Should UC1 teams (LIRR, OMB) require separate subscriptions for billing chargeback?	Pending finance/compliance review. Default: single sub with
OPEN Q	Is IaC (Bicep/Terraform) the agreed deployment method or will portal-based provisioning be allowed?	Recommendation: IaC-only via approved pipeline. Portal ch

ect
n tag-based cost allocation. Owner: Finance + Platform
hanges rejected by Policy. Owner: Platform Lead

🏗️ AI Foundry Platform – Hub-Based Architecture

A centralised Azure AI Foundry Hub hosting multiple agency projects with Hub-and-Spoke networking, RBAC isolation, and three onboarding patterns

● Status: DRAFT ● Azure AI Foundry ● Hub-Based Model ● Version 1.0 ● March 2026

📋 Table of Contents

1. Overview 2. Architecture Model 3. Network Topology 4. Use Cases 5. Azure Resource Inventory 6. Security Controls 7. RBAC Model 8. Naming Conv

1. Overview

The **AI Foundry Platform** is a **Hub-Based Azure AI Foundry deployment** that enables multiple agency teams to onboard AI projects quickly, securely, and consist

This architecture was designed to address the three most common onboarding patterns encountered across agency teams, ranging from fully private agent deploym

What this repo/document contains: Architecture diagrams (Pages 1–2), PRD & Technical Design (Pages 3–4), and this README (Page 5). The draw.io file is the

2. Architecture Model

This platform uses the **Hub + Projects model** — one AI Foundry Hub owns the shared infrastructure and acts as the governance boundary. Each team onboards as

Hub responsibilities: Owns the AI Foundry resource, shared connected resources (Azure OpenAI, AI Search, Storage, Key Vault), VNet peering configs, and Priva

Attribute	Hub-Based (THIS architecture)	Fou
Cost	✅ Lower – shared OpenAI, Search, Storage	❌ Higher – duplicated resources per team
Governance	✅ Centralised at Hub (policies, DNS, monitoring)	❌ Per-project, inconsistent across teams
Team isolation	✅ RBAC + subnet-level isolation per project	✅ Full resource-level isolation (overkill for m
Onboarding speed	✅ Fast – new Project under existing Hub (<5 days)	❌ Slow – new Hub + infrastructure per team

3. Network Topology

The platform uses an **Azure Hub-and-Spoke VNet topology** within a single subscription. All egress traffic is force-tunnelled through Azure Firewall Premium in the

Component	Address Space	Key Resourc
Hub VNet	10.0.0.0/16	Azure Firewall Premium, VPN/ExpressRoute GW, Bastion, Private DNS Zones, UDR, NSG
Spoke VNet – UC1 (Private Agent)	10.1.0.0/16	Agent Subnet (10.1.1.0/24) + Private Endpoint Subnet (10.1.2.0/24) | pe-llm, pe-data, pe-sr
Spoke VNet – UC2 (LLM + Search)	10.2.0.0/16	App Subnet (10.2.1.0/24) | Service Endpoints: Microsoft.CognitiveServices, Microsoft.Stora
Spoke VNet – UC3 (Shared Resources)	10.3.0.0/16	Shared Services Subnet (10.3.1.0/24) | Cross-sub Private Endpoints to Cosmos DB, Data

4. Use Cases

Use Case 1 – Private Agent Capability

Target teams: LIRR, OMB (high-security, regulated data environments)

Pattern: Fully private. Zero public internet path. All PaaS resources accessed through Private Endpoints only.
Network: Dedicated Agent Subnet + Private Endpoint Subnet (Subnet1). NSG deny-all inbound. UDR to Firewall.
Resources: Azure OpenAI (pe-llm), AI Search (pe-srch), Cosmos DB (pe-data), Azure Storage, Key Vault — all via PE.
DNS: Private DNS A-records in Hub DNS Zones (privatelink.openai.azure.com, privatelink.search.windows.net, privatelink.documents.azure.com).
When to use: PII data, regulated workloads, zero-trust requirement, agency mandates no public endpoint.

s.

ventions 9. Onboarding Guide 10. IaC & Deployment 11. Prerequisites 12. FAQ 13. Contact

tently. Rather than each team independently provisioning AI infrastructure, all projects share a single AI Foundry Hub — isolating teams through Azure RBAC and dedicated VN
nents requiring zero public internet exposure, to lightweight LLM access, to projects that consume shared enterprise data resources.
single source of truth for all design artifacts.

s a **Project** under this Hub. The Hub is *not* a separate deployment per team (that would be the Foundry Projects-only model, which is higher cost and loses shared governance)
te DNS Zones. **Project responsibilities:** Team RBAC, per-project connected resources (where UC requires), private endpoints scoped to the project spoke, and workload com

ndry Projects-Only (alternative)
ost)

Hub VNet. Each use case gets a dedicated Spoke VNet peered to the Hub.

es
s, DDoS Standard
rch | NSG: deny-all inbound
age, Microsoft.KeyVault | NSG: deny inbound
Lake, AI Search | Managed Identity access

let spoke networks.

).

npute.

Use Case 2 – LLM + AI Search Access Only

Target teams: Standard agency teams needing AI inference and search, no sensitive data

Pattern: Service Endpoints + VNet restrictions. No private endpoints required. Simpler, lower cost.
Network: App Subnet. Service Endpoints for CognitiveServices, Storage, KeyVault. NSG deny inbound, allow HTTPS to Azure service tags.
Resources: Azure OpenAI (VNet SR), AI Search (IP allowlist + AAD), Cosmos DB (Service Endpoint), Storage, Key Vault.
Monitoring: NSG flow logs + Log Analytics. All denied traffic captured.
When to use: Internal data, non-regulated, team primarily needs LLM calls and search indexing.

Use Case 3 – Shared Agent Resources

Target teams: Projects that consume shared enterprise data platforms

Pattern: Cross-subscription Private Endpoints + Managed Identity. Resources live outside AI Foundry in a shared services subscription.
Resources: Cosmos DB, Data Lake Storage Container, AI Search (all shared, PE from Spoke). Key Vault in shared sub holds connection strings.
Access control: Each project's Managed Identity is granted Reader/Data Reader role on shared resources. No wildcard grants. Read-only access enforced.
DNS: Private endpoint DNS A-records registered in Hub DNS Zones for cross-sub resolution.
When to use: Data already exists in enterprise data lake; duplication not acceptable; AI Foundry project reads from shared platform.

[redacted]

[redacted]

AI Foundry Platform – README (cont.) | Sections 5–13

5. Azure Resource Inventory

Resource	SKU / Tier	UC1	UC2	UC3	
AI Foundry Hub	Standard S0	Shared	Shared	Shared	Single Hub hosts all projects
Azure OpenAI Service	Standard S0	Private EP	Svc Endpoint	Shared	GPT-4o deployed per Hub
Azure AI Search	Standard S1	Private EP	IP Restrict	Shared index	Semantic search enabled
Cosmos DB	Serverless	Private EP	Svc Endpoint	Shared PE	NoSQL; auto-scale RU
Azure Storage (ADLS)	Standard LRS	Private EP	Svc Endpoint	Shared PE	Hierarchical namespace on
Azure Key Vault	Standard	Private EP	Svc Endpoint	Shared PE	Soft-delete + purge protect
Azure Firewall	Premium	Hub shared	Hub shared	Hub shared	IDPS enabled; TLS inspect
Log Analytics WS	Pay-as-you-go	All	All	All	Centralised; 90-day retain
VNet (Spoke)	N/A	10.1.0.0/16	10.2.0.0/16	10.3.0.0/16	Each peered to Hub VNet
Managed Identity	System-assigned	Per project	Per project	Per project	Used for all resource auth
Azure Bastion	Standard	Hub shared	Hub shared	Hub shared	Secure RDP/SSH; no public IP
Container Registry	Standard	Optional	Optional	Shared	For AI model container images

6. Security Controls

🔐 Identity & Authentication
All authentication via Azure Active Directory (Entra ID). No shared keys, passwords, or connection strings in code. Managed Identity per project for all resource access. Conditional A

🛡️ Network Security
NSG deny-all inbound on all subnets. Outbound restricted to Azure service tags only. Azure Firewall Premium with IDPS enabled. TLS inspection active. Force-tunnel via UDR – no d

🔑 Data Protection
All storage encrypted with Customer-Managed Keys (CMK) in centralised Key Vault. TLS 1.2+ enforced via Azure Policy. Key Vault soft-delete (90 days) and purge protection enable

📊 Monitoring & Alerting
All resources emit diagnostic logs to centralised Log Analytics Workspace. NSG flow logs and Azure Firewall logs enabled. Defender for Cloud Standard on subscription. Alerts for: fa

📋 Policy & Compliance
Azure Policy enforces: no public endpoints (UC1), TLS 1.2+, CMK encryption, approved regions only, required tags (env, team, project, cost-centre). Deny effect on non-compliant re

7. RBAC Model

Role	Assigned To	Scope	
Hub Owner	Platform Team only	AI Foundry Hub	Create/delete projects,
Project Contributor	Team AAD Group	AI Foundry Project (scoped)	Deploy models, manag
Reader	Stakeholders / Auditors	AI Foundry Project	Read-only view of proje
Managed Identity – Data Reader	Project Managed Identity	Shared Resources (UC3)	Read data from Cosmo
Key Vault Secrets User	Project Managed Identity	Key Vault (scoped)	Read secrets. Cannot s
Network Contributor	Platform Team only	Hub VNet + Spoke VNets	Manage peering, NSGs
Log Analytics Reader	All Project Teams	Log Analytics Workspace	Query own project logs

8. Naming Conventions

Resource Type	Naming Pattern	
Resource Group	rg-aif-{env}-{team}-{region}	rg-aif-prod-lirr-eus
AI Foundry Hub	aih-{org}-{env}-{region}	aih-nymta-prod-eus

Notes

ccess enforced for admin operations.
direct internet egress from any spoke.
d. No plaintext secrets in repos or configs.
ailed auth, unusual egress, policy violations, quota thresholds.
sources. Policy compliance dashboard reviewed monthly.

Permissions
manage Hub settings, assign project-level roles
e project resources, view logs. Cannot access sibling projects.
ct resources and outputs. No write access.
s DB, Data Lake, AI Search. No write, no delete.
set, delete, or purge secrets. Scoped per project vault.
s, UDRs, private endpoints. Not granted to project teams.
only (workspace scoped by resource tags).

Example

AI Foundry Project	aip-{team}-{env}	aip-lirr-prod
VNet (Hub)	vnet-hub-{org}-{env}	vnet-hub-nymta-prod
VNet (Spoke)	vnet-spoke-{team}-{env}	vnet-spoke-lirr-prod
Subnet (Agent)	snet-agent-{team}-{env}	snet-agent-lirr-prod
Subnet (PE)	snet-pe-{team}-{env}	snet-pe-lirr-prod
Private Endpoint	pe-{svc}-{team}-{env}	pe-llm-lirr-prod / pe-srch-lirr-prod
Key Vault	kv-{team}-{env}-{region}	kv-lirr-prod-eus (max 24 chars)
Log Analytics WS	law-aif-{org}-{env}	law-aif-nymta-prod
NSG	nsg-{subnet}-{team}-{env}	nsg-agent-lirr-prod
Managed Identity	id-{team}-{env}	id-lirr-prod

9. Onboarding Guide (Quick Start)

Step 1: Submit Onboarding Request
Open a service desk ticket. Provide: team name, project name, use case (UC1/UC2/UC3), data classification level, list of required resources, designated technical lead, and Azure AD

Step 2: Platform Team Review
Platform team validates use case selection, checks subnet CIDR availability, confirms resource requirements, and assigns VNet address space from the master IPAM table.

Step 3: Provision via IaC
Platform team runs the appropriate Bicep/Terraform module: • UC1: module/uc1-private-agent • UC2: module/uc2-llm-search • UC3: module/uc3-shared-resources All modules are pr

Step 4: Configure RBAC
Assign team AAD group to Contributor role scoped to their AI Foundry Project. Create system-assigned Managed Identity for the project. For UC3: assign Reader role on shared resc

Step 5: Connect Resources
Register connected resources in the AI Foundry Project. UC1: deploy private endpoints + DNS A-records in Hub zones. UC2: configure service endpoints + IP restrictions. UC3: regis

Step 6: Validate & Test
Platform team runs connectivity validation script. Verifies: DNS resolution, private endpoint reachability, RBAC isolation (no cross-project access), NSG rules, Firewall logs clean.

Step 7: Handoff & Go Live
Team receives: project URL, access credentials, runbook, and monitoring dashboard link. Project marked LIVE in the registry. Onboarding ticket closed.

10. IaC & Deployment

Tooling: Bicep (primary) + Terraform modules (optional). All deployments via Azure DevOps pipeline or GitHub Actions. No manual portal deployments — enforced by Azure Policy (

Repo structure: /infra/hub/ (Hub VNet, Firewall, Bastion, DNS) | /infra/modules/uc1/ | /infra/modules/uc2/ | /infra/modules/uc3/ | /infra/shared/ (Cosmos DB, Data Lake, Key Vault) | /s

Pipeline flow: lint → plan → security scan (Checkov/tfsec) → approval gate → apply → smoke test → notify team

State management: Terraform state in Azure Blob Storage (private endpoint, versioning enabled, RBAC restricted to pipeline SP only).

11. Prerequisites

Azure Subscription	Active subscription with Owner or Contributor rights. Quota pre-increased for OpenAI (GPT-4o), AI Search, Cosmos DB in target reg
Azure AD / Entra ID	Global Admin or Privileged Role Administrator to assign RBAC roles. AAD groups created per team before onboarding.
Networking	IP Address Management (IPAM) plan agreed. Hub VNet deployed (or deploy via /infra/hub/). ExpressRoute or VPN if on-prem conne
Tooling	Azure CLI 2.55+, Bicep CLI 0.23+, Terraform 1.6+ (optional), Python 3.10+ (for validation scripts), Git access to this repo.
AI Foundry Permissions	Subscription-level permission to create AI Foundry Hub resource. Microsoft.MachineLearningServices provider registered on subscr
Defender for Cloud	Defender for Cloud Standard plan enabled on subscription BEFORE first project onboards (required for UC1 compliance).

12. Frequently Asked Questions

Q: Do the 3 use cases require separate subscriptions?
A: No. All three use cases run within a single subscription. Isolation is achieved through RBAC (scoped per project) and dedicated VNet spokes. Separate subscriptions are only nee

Q: Can UC1 and UC2 teams share an AI Search index?
A: No by default. Each project has its own connected AI Search resource. UC3 is the pattern for intentional resource sharing — it provides controlled read-only access via Managed I

D group name.

re-approved and policy-compliant.

ources.

ster cross-sub private endpoints.

deny deployments outside approved pipeline service principal).

scripts/validate-onboarding.sh

gion (e.g. East US).

ectivity is needed for UC1.

ription.

ded if there is a hard compliance boundary (e.g. FedRAMP vs standard) or a billing chargeback requirement that cannot be satisfied with cost tags.

Identity with explicit RBAC grants.

Q: What happens if a team needs both private endpoints AND shared resources?
A: A project can be a hybrid of UC1 and UC3. The private endpoint pattern from UC1 is applied for connectivity, and the shared resource access model from UC3 is applied for the ex

Q: Can teams deploy their own resources outside the AI Foundry Project?
A: Only within their scoped Resource Group. Azure Policy prevents deploying public-facing resources or resources in non-approved regions. All networking changes must go through

Q: How are LLM model deployments managed?
A: Model deployments (e.g. GPT-4o, Ada-002) are managed at the Hub level by the Platform Team. Teams request new deployments via ticket. Quota is allocated per project using A

Q: Is the Hub a single point of failure?
A: The AI Foundry Hub is a control plane resource. It does not process inference traffic — that goes directly to Azure OpenAI. Hub unavailability affects project management (portal/A

13. Contributing & Contact

Document Owner: AI Platform Team — responsible for all architecture decisions, IaC modules, and onboarding approvals.

To request changes to this architecture: Raise a PR against this draw.io file in the platform repo, or open an Architecture Review Board (ARB) ticket for major changes (new use c

To onboard a new project: Follow Section 9 — submit a service desk ticket with the required fields. SLA: 5 business days from approved ticket to live project.

Security incidents: Report immediately to the Security Operations Centre (SOC). Do not raise a standard service desk ticket for suspected breaches.

xternal Cosmos DB / Data Lake. Document this as a custom variant in the onboarding ticket.

n the Platform Team.

Azure OpenAI deployment quotas.

API), not live AI inference calls already in flight.

case, subscription changes, security control changes).

Team